

LUIS ANTONIO REYNOSO SALCEDO

Lead AI Engineer | Forward Deployed Engineer | LLM Agents · RAG · MCP | MLOps & Cloud Platform

Guadalajara, Mexico | luisantonioareysal@gmail.com | linkedin.com/in/areynoso | github.com/eltelemaco

EXECUTIVE SUMMARY

AI Engineer and Forward Deployed Engineer with **15+ years** of experience building production systems across **AWS, Azure, and GCP** for Fortune 500 enterprises. Designs and ships **tool-using LLM agents, RAG assistants, and multi-model AI platforms** with **Claude Code, MCP, LiteLLM, Amazon Bedrock, SageMaker, Vertex AI, and Azure AI Foundry**, and builds the **evaluation, guardrails, and observability** that keep them reliable in production. Works directly with customers and business stakeholders to turn their workflows into deployed AI solutions. Hands-on field experience includes integrating in-vehicle AI for **2M+ vehicles** (Daimler) and translating customer requirements into API, networking, and observability specifications. Strong platform foundation: managed **3,000+ microservices** across **6+ global regions** at **99.99% uptime**, cut deployment cycles by **75%**, automated **90%** of infrastructure provisioning, and delivered **\$1M+** in annual cost savings for systems handling **100M+ daily transactions**.

KEY ACHIEVEMENTS

- **Agentic Engineering Platform (KKR, Forward Deployed):** Built tool-using LLM agents on Claude Code with custom MCP servers that connect them to GitHub, Terraform, AWS, and observability APIs. The agents fix CI/CD failures, propose infrastructure changes, and generate runbooks, with every change approved by a human through a pull request. They run on a LiteLLM model layer spanning Bedrock, SageMaker, Vertex AI, and Azure AI Foundry, which significantly reduced manual cycle time for DevOps and ML teams.
- **Speech Recognition AI Cloud Migration (Daimler Automotive Group, Germany):** Led the migration of 20+ mission-critical speech-recognition AI microservices from on-premises to Azure Kubernetes Service, enabling real-time AI for 2M+ vehicles with zero service disruption. Worked hands-on with in-vehicle hardware and electronics integration and served as the direct customer-facing technical contact for Daimler engineering. Achieved 85% faster deployments, 40% lower latency, and 99.98% uptime.
- **AWS Landing Zone & ML Workload Migration (Global Renewable Energy Company):** Designed an AWS landing zone and migrated 15+ ETL pipelines plus the energy-forecasting and machine-learning applications. Enabled 300% burst scalability, cut infrastructure cost by 45%, and improved query performance by 60%.
- **Enterprise Landing Zone Architecture (Government Energy Regulatory Entity):** Architected a greenfield multi-tenant, multi-region Azure landing zone aligned with SOC 2 and FedRAMP. Reached 95% Terraform coverage, 3x faster provisioning, and a 99.95% RPO/RTO SLA.

CORE TECHNICAL COMPETENCIES

- **AI Engineering & LLM Agents:** AI Agents & Tool Calling | Agentic Workflows | Model Context Protocol (MCP) Servers | Claude Code | RAG (chunking, embeddings, vector search) | Prompt Engineering & Structured Outputs | Conversational AI / Chatbots | AI-Assisted Software Development | Coding Agent Orchestration
- **LLM Platforms & Models:** Amazon Bedrock | Amazon SageMaker | Google Vertex AI | Azure AI Foundry | Azure OpenAI | LiteLLM (multi-model gateway & routing) | Claude / GPT / Gemini APIs | GitHub Copilot | Cursor
- **AgentOps, Evaluation & MLOps:** LLM Evaluation & Regression Testing | Agent Tracing & Token/Cost Monitoring | Guardrails, Sandboxing & Least-Privilege Tooling | Human-in-the-Loop Automation | Model Deployment & Serving | AI Inference | ML Pipeline Automation | Model Monitoring | CI/CD for ML
- **APIs & Service Delivery:** Python | FastAPI | REST | gRPC | Swagger/OpenAPI | Software Delivery Lifecycle (SDL) Automation
- **Cloud Platforms:** AWS (EKS, ECS, Lambda, S3, RDS, Aurora, Redshift, Athena, KMS, IAM, CloudWatch) | Azure (AKS, Functions, App Services, Azure DevOps) | GCP
- **Infrastructure & Containers:** Terraform | Ansible | CloudFormation | Bicep | ARM | Docker | Kubernetes (EKS, AKS, GKE) | Helm | Coder Remote Dev Environments
- **CI/CD & Release Engineering:** GitHub Actions | GitLab CI/CD | Jenkins | Azure DevOps | CircleCI | CodePipeline | Liquibase
- **Networking & Field Deployment:** Site-to-Site VPN & VPC Peering | TLS/PKI | Firewalls / NSGs / Security Groups | Wi-Fi & LTE/5G | In-Vehicle / Edge Hardware Integration | IoT & Sensor Integration | Customer-Facing Field Deployment

- **Data Platforms:** Databricks | Apache Spark | Snowflake | PostgreSQL | Airflow | Dagster | AWS Glue | AWS Batch | Azure Data Factory | GCP Data Fusion
- **Observability & SRE:** OpenTelemetry | Datadog | Splunk | Dynatrace | ELK | CloudWatch | PagerDuty | SLO/SLI Definition | Incident Response | Root Cause Analysis
- **Languages:** Python | Go | Bash | PowerShell | SQL | YAML | HCL

PROFESSIONAL EXPERIENCE

Lead AI Engineer: LLM Agents & Platform (Forward Deployed) *Movate / KKR | Boston, MA*

March 2026 – August 2026

Official title: Senior DevOps MLOps Data Platform Engineer

- **Agentic Systems Engineering:** Designed and built tool-using LLM agents in Python on Claude Code and the Model Context Protocol (MCP). Wrote custom MCP servers that expose GitHub, Terraform, AWS, and observability APIs as agent tools, so agents can fix CI/CD failures, propose infrastructure changes, and generate runbooks from start to finish. Every action goes through a pull request for human approval.
- **Multi-Model LLM Architecture:** Built a LiteLLM-based model layer that gives teams one interface to Amazon Bedrock, SageMaker endpoints, Google Vertex AI, and Azure AI Foundry, with routing by task, fallback across providers, and centralized tracking of token usage and cost.
- **Conversational AI & RAG:** Built an internal engineering chatbot that answers questions from runbooks, pipeline documentation, and platform knowledge using retrieval-augmented generation, cutting the time engineers spend searching for operational answers.
- **Customer-Facing AI Delivery:** Embedded with insurance and data-platform stakeholders to scope AI use cases. Deployed production inference services on Bedrock, Vertex AI, and Azure AI Foundry that power real-time forecasting, risk, and fraud-analysis features in insurance agent applications, translating each workflow into prompts, APIs (FastAPI/REST/gRPC), and success metrics.
- **LLM Evaluation & AgentOps:** Set up the reliability practices for production agents: versioned prompts, regression tests on curated task sets, checks of tool-call success and output quality, and tracing of agent runs, latency, and token spend. Agents ran in sandboxed Coder workspaces with least-privilege IAM and tool allow-lists.
- **Production AI Platform Foundations:** Built greenfield GitHub Actions CI/CD for the AI and data platform, plus Liquibase-driven, environment-aware Redshift schema releases through formal change management. Hardened the Terraform/IAM setup with tighter KMS policies, cross-account SNS patterns, and centralized Lambda execution roles across four workspaces.

Senior AI / MLOps Engineer: AI Platform & Intelligent Automation *EPAM Systems | Houston, TX*

May 2023 – Oct 2025

Official title: Senior Release DevOps Engineer / SRE

- **LLM-Assisted Incident Intelligence (AIOps):** Built an LLM-assisted incident workflow that summarizes alerts from Splunk and Dynatrace, correlates related signals, drafts root-cause hypotheses and post-mortems, and suggests runbook steps. Combined with automated alerting, it cut MTTR from 90 to 15 minutes across 200+ incidents and reduced repeat incidents by 60%.
- **Generative AI Developer Enablement:** Rolled out AI-assisted development (GitHub Copilot and LLM-generated templates for Terraform modules and CI/CD pipelines) to engineering teams, speeding up service onboarding and contributing to a 90% reduction in manual provisioning across 500+ resources.
- **Model & Service Delivery Platform:** Built CI/CD for models and services, with packaging, artifact promotion, automated validation gates, and progressive rollout on EKS/AKS. Cut deployment time from 25 to 4 minutes (82%) and the failure rate from 8% to 0.3%.
- **Scalable AI/ML Infrastructure:** Ran 100+ Kubernetes clusters (2,000+ nodes) hosting 2,500+ microservices and model-serving workloads at 50M+ daily transactions and a 99.99% uptime SLA. IaC provisioning cut cluster setup time from 2 hours to 12 minutes.
- **Compute Cost Optimization:** Saved \$1.2M a year (35%) through data-driven right-sizing, reserved-instance strategy, and autoscaling tuning across application and ML compute, and cut wasted compute by 40%.
- **Enterprise Cloud Migration:** Led a Fortune 500 migration of 1,800+ on-premises servers to AWS that finished 3 months early with zero data loss and no downtime, and set up the cloud foundation for later AI and data workloads.

Senior DevOps Engineer / Release Engineer *Tecnologia Corp | Houston, TX*

November 2021 – May 2023

- **Kubernetes Scaling & Automation:** Architected AKS infrastructure running 3,000+ containerized applications across 5 regions; automated cluster provisioning with Helm and CLI tooling, cutting deployment time by 75%.
- **Internal Platform Tooling:** Designed a proprietary self-service provisioning toolset for 100+ development teams, cutting infrastructure request fulfillment from 3 days to 2 hours.
- **CI/CD Pipeline Enhancement:** Optimized Azure DevOps and AWS CodePipeline, cutting build-to-deploy time from 60 to 15 minutes, and added parallel testing that raised test coverage to 92%.
- **Team Leadership & Mentorship:** Led a DevOps team of 8 engineers and mentored 12+ junior engineers; process improvements raised team productivity by 40%.
- **Knowledge Base & Runbooks:** Built a runbook library covering 200+ operational procedures; the self-service knowledge base cut support tickets by 45%.
- **Cloud Migration Strategy:** Designed cloud migration roadmaps with enterprise stakeholders for 2+ Fortune 500 companies moving to Azure and AWS, achieving 85% cost reduction versus on-premises.

Release Engineer / DevOps SRE Tech Lead *Luxoft (DXC Technology) | San Francisco, CA*

June 2017 – Nov 2021

- **In-Vehicle AI Deployment (Daimler "Hey Mercedes", 2020):** Worked hands-on with in-vehicle hardware and electronics integration for a custom speech-recognition model trained in Azure and served through a global API to 2M+ vehicles in real time. Served as the direct customer-facing technical contact, coordinating hardware and software requirements between Daimler engineering and the delivery team.
- **Microservices Deployment at Enterprise Scale:** Ran continuous deployment of 4,000+ microservices across travel, automotive, healthcare, and financial services, with a 99.97% deployment success rate.
- **Observability Architecture:** Architected an enterprise Splunk stack processing 10TB+ of logs daily from 5,000+ servers, plus Dynatrace APM dashboards, cutting incident discovery time from 4 hours to 8 minutes.
- **Incident Response Excellence:** Led incident response for a critical production outage and used automation to cut MTTR from 120+ to 18 minutes, maintaining 99.95% availability.
- **Automation at Scale:** Wrote 50+ Python and Bash automation scripts monitoring 8,000+ application instances, cutting manual notifications by 85%.
- **CI/CD Pipeline Leadership:** Improved CI/CD on Jenkins, CircleCI, and GitHub Actions for 200+ development teams, cutting code-to-production time by 65%.

DevOps Engineer / SRE Tech Lead *Accenture | Monterrey, NL*

June 2016 – June 2017

Cloud Engineer / System Administrator *Softtek | Aguascalientes*

Oct 2015 – June 2016

- **Warehouse Automation & Logistics Systems:** Built and supported inventory-tracking and logistics automation across multiple warehouses for a large beverage-industry client, connecting IT systems to warehouse operations for real-time inventory control.

Systems Engineer / Server Administrator *TATA TCS | Guadalajara*

Oct 2012 – Oct 2015

Systems Server Administrator *City Council IT Governance | Jalisco*

May 2009 – Sept 2012

CERTIFICATIONS & PUBLICATIONS

- Splunk Administrator Certified (Cloud, Cluster Administration, Enterprise Security, IT Service Intelligence, Big Data Analytics)
- ITIL Foundations V3 Certified Professional
- International Program in Cybersecurity and Big Data, Barcelona, Spain (2019)
- Data Science Oriented Programming and Machine Learning, Guadalajara, Mexico (2018)
- *Publication:* Velarde Peña, L. & Reynoso Salcedo, L. A. "Robotic Arm D-Tran RT3200 as Educational Prototype." First International Congress of Trans-disciplinarity, Mexicali, BC, Mexico (2010). ISBN: 978-607-7753-55-1

EDUCATION

Bachelor of Science in Computer Systems Engineering · Instituto Tecnológico Superior de Puerto Vallarta, México (2009)